From Awareness to Institutionalization:

A Change Management Framework for Organizational AI Epistemic Governance

Paper 10 in the AI Epistemic Governance Stack — The Deployment Layer for Papers 7–9

Series DOIs: 10.5281/zenodo.19562751 (Paper 7) | 10.5281/zenodo.19565086 (Paper 8)

Paper 9 (ZTEK) and Paper 10 (this paper) submitted concurrently to Zenodo

Executive Brief

This brief is intended for senior leaders and decision-makers who need the key points of this paper in two minutes.

The problem in one sentence:

Your organization is almost certainly using AI-generated outputs without systematic verification — and the errors this produces are indistinguishable from correct results until they surface downstream.

What this paper provides:

- A five-phase roadmap for moving from unmanaged AI usage to institutional epistemic governance
- The Epistemic Risk Map: a one-page visual that shows where your organization is most exposed
- A complete resistance map: 8 common objections and evidence-based responses
- A pilot design that produces measurable results before full adoption is required

Three measurable outcomes from pilot implementations:

Metric	Before Brakes	After Brakes (ZTEK Pilot)	Change
Unverified claims entering knowledge base	~100% of AI outputs	~30% of AI outputs (those passing reproduction and/or verification — T1 or above)	~70% reduction
Fabricated claims caught before propagation	0 (no gate)	3 caught in 8-week pilot	Systematic detection
Verification overhead per claim cycle	0 (no process)	~20-30% additional time	Accepted as worthwhile

The decision this paper asks of you:

Not full adoption. Start with Phase 1 (Awareness) and Phase 2 (Assessment). The Epistemic Risk Map alone — a single deliverable — gives your organization an actionable picture of AI-related epistemic exposure with no technical prerequisites.

Series context:

This is Paper 10 in a four-paper AI Epistemic Governance Stack. Paper 7 explains why AI fails. Paper 8 tracks failure over time. Paper 9 (ZTEK) enforces the boundary. This paper makes the stack deployable at organizational scale.

Abstract

Technical adoption of AI governance frameworks is relatively straightforward. Organizational adoption is hard. An organization can implement the ZTEK protocol technically in a week — but if the people using it do not understand why, if leadership does not mandate it, and if incentives still reward output speed over verification discipline, the framework fails silently. The architectural brakes get bypassed not by the AI, but by the humans managing it.

This paper introduces a five-phase Change Management Framework for organizational adoption of the AI Epistemic Governance Stack (Papers 7–9). The framework addresses the human, structural, and cultural barriers that determine whether AI governance becomes embedded practice or shelfware. It is designed for CIOs, AI governance leads, research directors, compliance teams, and any organizational leader responsible for AI deployment at scale.

The framework is structured around five phases: Awareness, Assessment, Pilot, Institutionalization, and Continuous Improvement. Each phase has defined deliverables, resistance patterns, and success criteria. A pilot vignette from the reference implementation (PCF research program) provides concrete before/after metrics. The Epistemic Risk Map is provided as a standalone tool usable without prior framework adoption.

Table of Contents

Executive Brief

1. Introduction: Why Technical Adoption Is Not Enough

2. The Complete AI Epistemic Governance Stack

3. The Five-Phase Change Management Framework

4. Phase 1 — Awareness

5. Phase 2 — Assessment

6. Phase 3 — Pilot

7. Phase 4 — Institutionalization

8. Phase 5 — Continuous Improvement

9. Organizational Resistance: The Full Resistance Map

10. Roles and Responsibilities

11. Metrics and Success Criteria

12. The Epistemic Risk Map: A Standalone Tool

13. Case Study: Research Organization Reference Implementation

14. Pilot Vignette: Before and After Metrics

15. Adoption Roadmap

16. Conclusion

References

1. Introduction: Why Technical Adoption Is Not Enough

The history of enterprise technology is full of frameworks that were technically sound and organizationally ignored. The reasons are consistent: the framework required behavior change, the incentives did not support the change, and no one made the change visible enough to sustain.

AI governance is following the same pattern. Organizations adopt AI tools rapidly. They adopt AI governance slowly, reluctantly, and often only after a visible failure.

This paper is written for the period before that failure — for organizations that want to build epistemic discipline into their AI workflows before they need to explain to stakeholders why an AI-generated result was wrong, widely cited, and undetectable until it was too late.

1.1 The Core Problem This Framework Addresses

Papers 7, 8, and 9 in this series establish that:

- AI systems lack the internal mechanisms humans rely on to stop themselves from fabricating (Paper 7)
- An accountability layer is needed to track and score AI reliability over time (Paper 8: AEAL)
- A boundary enforcement protocol is needed to control what enters the verified knowledge base (Paper 9: ZTEK)

Paper 10 addresses the question those three papers leave open: how does an organization actually adopt this?

The answer is not a technical implementation guide. It is a change management framework — a structured pathway for moving an organization from unmanaged AI usage to institutionalized epistemic governance.

Throughout this paper, the term brakes refers to the architectural controls introduced in Paper 9 (ZTEK) that prevent unverified AI output from entering the knowledge base. These controls are structural, not behavioral — they enforce verification at the boundary regardless of the fluency or apparent confidence of the AI output.

1.2 Who This Framework Is For

Role	What They Need From This Framework
CIO / CTO	Strategic roadmap and business case for AI governance investment
AI Governance Lead	Operational framework for policy design and enforcement
Research Director	Workflow integration guide for AI-assisted knowledge work

Role	What They Need From This Framework
Compliance / Legal	Risk exposure map and audit trail requirements
Team Leads / Managers	Practical phase-by-phase adoption guide
Individual Contributors	Clear explanation of why the framework exists and what it requires of them

2. The Complete AI Epistemic Governance Stack

Paper 10 is the deployment layer of a four-paper stack. Understanding the full stack is necessary before adopting any part of it.

Paper	Layer	Core Question Answered
7: AI Behavior Science	Theoretical Foundation	Why does AI fail epistemically? (DOI: 10.5281/zenodo.19562751)
8: AEAL Framework	Accountability Layer	How do we track and score AI reliability? (DOI: 10.5281/zenodo.19565086)
9: ZTEK Protocol	Enforcement Mechanism	How do we enforce the boundary between AI output and verified knowledge?
10: Change Management (this paper)	Deployment Layer	How does an organization actually adopt this?

Organizations do not need to adopt all four layers simultaneously. The framework is designed for incremental adoption starting from Phase 1 (Awareness) with zero technical prerequisites.

3. The Five-Phase Change Management Framework

The framework moves an organization through five phases, each building on the last. The phases are not strictly sequential — organizations may run Phase 2 and Phase 3 in parallel, or revisit Phase 1 after a Phase 3 pilot reveals new gaps. The sequence is a guide, not a constraint.

Phase	Name	Core Activity	Primary Deliverable
1	Awareness	Build organizational understanding of AI epistemic risk	Leadership brief, literacy baseline
2	Assessment	Map current epistemic exposure	Epistemic Risk Map
3	Pilot	Test the framework in a	Pilot results, validated

Phase	Name	Core Activity	Primary Deliverable
		controlled workflow	brake system
4	Institutionalization	Embed the framework into policy and incentives	AI Governance Policy, trust tier standards
5	Continuous Improvement	Evolve the framework with the AI landscape	Quarterly audit, failure mode registry

4. Phase 1 — Awareness

The organization does not yet know what it does not know. Phase 1 builds the foundational understanding that makes everything else possible.

4.1 Objective

Move key stakeholders from "our AI is probably fine" to "our AI is probably fabricating something right now, and we don't know what." This is not alarmism — it is the accurate baseline from which governance becomes rational.

4.2 Key Interventions

- **AI Behavior Science literacy brief:** A one-page, non-technical summary of the core structural differences between humans and AI (from Paper 7). Designed for leadership, not engineers.
- **Fabrication assumption shift:** A facilitated discussion exercise in which teams identify one AI output from the past month that was used without independent verification. In most organizations, this exercise surfaces multiple examples immediately.
- **Visible failure audit:** Identify any past AI-assisted outputs that later turned out to be incorrect. Trace back to the point where the error entered the workflow. Document the path.

4.3 Primary Resistance Pattern

"Our use case is different."

Every organization believes its AI usage is lower-risk than it is. The fabrication assumption shift exercise is the most effective response — not argument, but demonstration. When teams identify their own recent unverified outputs, the resistance typically softens.

4.4 Success Criteria

- Leadership can articulate in plain language why AI fabrication is a structural risk, not a model quality problem
- At least one visible failure has been traced and documented
- The organization has committed to conducting a Phase 2 assessment

5. Phase 2 — Assessment

Phase 2 measures the organization's current epistemic exposure. The primary deliverable is the Epistemic Risk Map — a visual representation of every point where unverified AI output touches decisions, documents, or downstream systems.

5.1 Objective

Produce a complete picture of where AI output currently enters the organization's knowledge base, at what trust level it is treated, and what the consequences of a fabrication at each point would be.

5.2 Key Interventions

- **Workflow mapping:** For each AI-assisted workflow, identify: (a) where AI output is generated, (b) where it enters a document, decision, or downstream system, (c) whether any verification step exists between generation and entry.
- **Trust tier audit:** Classify each AI-output entry point by the ZTEK trust tier it is currently treated as versus the tier it has actually earned. The gap between these two is the epistemic risk exposure.
- **Relay chain identification:** Identify any workflows where AI output from one model or session is fed as input to another without a boundary verification step. These are highest-risk propagation paths.

5.3 Success Criteria

- Epistemic Risk Map completed and reviewed by leadership
- At least three high-risk entry points identified
- Phase 3 pilot workflow selected based on map findings

6. Phase 3 — Pilot

Phase 3 tests the framework in a controlled, single-workflow context. The goal is not comprehensive adoption — it is producing visible, measurable evidence that the brakes work.

6.1 Selecting the Pilot Workflow

The pilot workflow should satisfy four criteria:

- It has clear, verifiable inputs and outputs
- It currently involves AI generation without systematic verification
- A fabrication in this workflow would have detectable consequences
- The team running it is willing to participate in honest documentation

6.2 Running the Pilot

Baseline run (without brakes):

- Run the AI workflow normally
- Document all AI outputs
- Record: fabricated details, unverified claims, missing provenance, confident assertions without evidence

Brake-enabled run (with ZTEK):

- Apply the minimum viable ZTEK implementation: reproduce commands, evidence classes, boundary verification
- Document what changes: what gets caught, what gets rejected, what forces re-derivation
- Record: fabrication rate change, verification coverage, trust tier distribution

6.3 Measuring Pilot Results

Metric	Baseline	Target (Brake-Enabled)	Measurement Method
Fabrication Rate (FR)	Measure and record	Reduction vs. baseline	Independent claim audit
Verification Coverage	Measure and record	Increase vs. baseline	Provenance field completion
Trust Tier Distribution	Measure and record	T0 reduction, T2 increase	Tier registry count
Capability Honesty (CH)	Measure and record	Increase vs. baseline	AEAL CH axis scoring

6.4 Success Criteria

- Measurable reduction in unverified claims entering the knowledge base

- At least one fabricated claim caught by the boundary verification step that would previously have propagated
- Team reports that the process is sustainable
- Leadership commits to Phase 4 based on pilot results

7. Phase 4 — Institutionalization

Phase 4 embeds the framework into organizational policy, process, and incentives. This is where the framework moves from pilot practice to institutional norm.

7.1 Objective

Make ZTEK the default, not the exception. This requires three simultaneous changes: policy, process, and incentives. Changing only one or two of the three is the most common reason institutionalization fails.

7.2 Policy Changes

- ZTEK trust tier requirements encoded into AI usage policy: "No T0 output may be cited in any client-facing document, published paper, or organizational decision." (T0 = raw AI output, unverified.)
- Evidence class vocabulary adopted as standard terminology in AI-assisted workflows
- Relay chain verification requirements added to multi-model or multi-session workflows
- AI disclosure statements required on all outputs with significant AI contribution

7.3 Process Changes

- Provenance fields added to document templates for AI-assisted work
- Boundary verification step added to existing review and approval workflows
- Trust tier assignment added to project management and knowledge management systems
- Model upgrade protocol: when a new AI model is adopted, all outputs from that model re-enter at T1

7.4 Incentive Realignment

This is the most commonly neglected and most important component. If speed of output is still the primary reward signal, verification discipline will erode under pressure.

- Verification thoroughness added as an explicit performance criterion
- T3-certified outputs recognized and rewarded differently from T0 outputs
- Fabrication detection rewarded, not penalized — catching errors is a contribution, not an accusation

7.5 Operational Cost and ROI

A common concern at Phase 4 is operational cost. The following framework helps quantify the trade-off:

Cost Component	Estimate	Notes
Verification overhead per claim (T0 to T2)	15–45 min	Majority is verification, not paperwork
Policy and process change (one-time)	2–4 weeks	Scales with org size
Training per team member	4–8 hours	Amortized across all future AI-assisted work
Fabrication recovery cost (if undetected)	Days to weeks	Includes rework, reputational, and decision costs
Break-even point	~3–5 caught fabrications	Each caught fabrication avoids downstream recovery cost

7.6 Primary Resistance Pattern

"This slows us down."

The answer is not to deny the overhead — verification does add time. The answer is to make the alternative visible: undetected fabrication is slower. It bills the cost later, at the point where the error surfaces. The pilot results from Phase 3 are the most effective response to this resistance.

7.7 Success Criteria

- AI governance policy published and acknowledged by all AI-using teams
- Trust tier requirements embedded in at least three core workflows
- Incentive changes implemented and communicated
- First quarterly audit scheduled (Phase 5 trigger)

8. Phase 5 — Continuous Improvement

Phase 5 is not a destination — it is the ongoing operating mode of a mature AI governance organization. The AI landscape changes rapidly; the governance framework must evolve with it.

8.1 Quarterly Epistemic Audits

- Review trust tier distribution across all AI-assisted workflows
- Identify new high-risk entry points created by new AI tool adoption
- Review the failure mode registry for new incident patterns
- Update the Epistemic Risk Map

8.2 Failure Mode Registry

Every fabrication incident — caught or uncaught — is documented in the organizational failure mode registry. The registry serves three functions:

- Organizational learning: patterns across incidents inform protocol updates
- Audit evidence: demonstrates active governance to regulators and stakeholders
- Benchmarking: enables comparison across teams and over time

8.3 Model Upgrade Protocol

When a new AI model is adopted, the following steps are mandatory:

1. All outputs from the new model begin at T1 regardless of prior model trust history
2. A structured evaluation period is run before T2 promotion is permitted
3. The Epistemic Risk Map is updated to reflect the new model's deployment points
4. Team literacy brief updated to reflect any new capability or failure mode characteristics

9. Organizational Resistance: The Full Resistance Map

Resistance to AI governance frameworks is predictable. The following table documents the most common resistance patterns, their root causes, and the most effective responses.

Resistance Pattern	Root Cause	Most Effective Response
"Our use case is different"	Belief that AI risk is a general problem, not a specific one	Run the fabrication assumption shift exercise — demonstrate with their own recent outputs, not abstract examples
"This is too complex"	Overwhelm from full framework scope	Start with the Epistemic Risk Map only. One deliverable, no technical prerequisites, immediate

Resistance Pattern	Root Cause	Most Effective Response
		value
"We don't have time to verify everything"	Misunderstanding of ZTEK scope	Clarify: ZTEK is a gate at the boundary, not a blanket audit of everything. Only claims entering the knowledge base need verification
"This slows us down"	Accurate but incomplete cost accounting	Make the downstream cost visible. Show the pilot vignette metrics: 3 fabrications caught that would each have required days of recovery
"Our people will resist the process"	Anticipatory resistance, often inaccurate	Run the pilot first. Visible results change behavior faster than top-down mandates. Let the team report their own experience
"This only applies to research"	Narrow framing of AI knowledge work	Any decision informed by AI output is knowledge work. Legal, finance, HR, and strategy teams are all in scope
"The AI has gotten better, we don't need this"	Model capability conflated with epistemic reliability	Better models fabricate more fluently, not less. Capability and reliability are orthogonal dimensions
"We already have AI ethics policies"	Conflation of ethics policy with epistemic governance	Ethics addresses intent and values. ZTEK addresses reliability and verification. Both are needed; neither substitutes for the other

10. Roles and Responsibilities

Role	Phase 1	Phase 2	Phase 3	Phase 4	Phase 5
Executive Sponsor	Mandate awareness	Approve assessment	Review pilot results	Approve policy	Receive audit reports
AI Governance Lead	Run literacy brief	Lead assessment	Design and run pilot	Draft policy	Own audit process
Team Leads	Participate in brief	Map their workflows	Run pilot workflow	Implement process changes	Submit incident reports
Individual Contributors	Complete literacy brief	Provide workflow data	Run brake-enabled workflow	Follow new process	Document incidents
IT / Engineering	Not required	Support relay chain identification	Implement ZTEK gate	Maintain tier registry	Maintain systems

11. Metrics and Success Criteria

Phase	Key Metric	Target
1: Awareness	Leadership literacy score (0–5)	3+ before proceeding to Phase 2
2: Assessment	Epistemic Risk Map completion	100% of AI-assisted workflows mapped
3: Pilot	Fabrication rate reduction vs. baseline	Measurable reduction; at least one caught fabrication
4: Institutionalization	Policy compliance rate	90%+ of AI-assisted outputs meeting trust tier requirements within 6 months
5: Continuous Improvement	Quarterly audit completion rate	100%; failure mode registry maintained with every incident

11.1 Organization-Level Health Indicators

- Trust tier distribution trend: T2 and T3 share increasing over time
- Fabrication detection rate: incidents caught before propagation vs. after
- Policy compliance rate: outputs meeting provenance requirements
- Model upgrade protocol adherence: new models starting at T1

12. The Epistemic Risk Map: A Standalone Tool

The Epistemic Risk Map is the most immediately actionable tool in this framework. It can be produced without any prior framework adoption and used independently to prioritize governance investments.

12.1 Construction

To construct the Epistemic Risk Map, work through the following steps in order:

1. List all workflows that currently use AI-generated output in any form.
2. For each workflow, identify the specific entry point where AI output enters a document, decision, or downstream system.
3. Assign the current assumed trust tier at each entry point (T0–T3 using ZTEK definitions).
4. Assign the earned trust tier at each entry point based on actual verification evidence available.
5. Calculate the trust gap: assumed tier minus earned tier. A gap of 2 or more is high-risk.
6. Assess consequence severity: what is the downstream impact if a fabrication at this point propagates undetected? Rate as low / medium / high / critical.
7. Plot all entry points on the 2x2 matrix in Section 12.2: trust gap on the x-axis, consequence severity on the y-axis.

12.2 The Epistemic Risk Map Matrix

Figure 2. Epistemic Risk Map — 2x2 priority matrix (trust gap x consequence severity). Plot each AI-output entry point as a node. Upper-left quadrant = critical priority.

HIGH CONSEQUENCE	HIGH CONSEQUENCE
HIGH TRUST GAP HIGH CONSEQUENCE **CRITICAL PRIORITY** Implement ZTEK immediately	LOW TRUST GAP HIGH CONSEQUENCE **HIGH PRIORITY** Verify existing controls
HIGH TRUST GAP LOW CONSEQUENCE **MEDIUM PRIORITY** Scheduled remediation	LOW TRUST GAP LOW CONSEQUENCE **MONITOR** Quarterly review sufficient
HIGH TRUST GAP ← x-axis → LOW TRUST GAP	

12.3 Interpretation

The upper-left quadrant (high trust gap, high consequence severity) represents the highest-priority governance targets. These are the workflows where the organization is most exposed and where ZTEK implementation will produce the greatest risk reduction per unit of effort. Start Phase 3 pilots here.

The lower-right quadrant (low trust gap, low consequence) represents workflows where monitoring is sufficient. Do not over-invest governance effort in this quadrant.

13. Case Study: Research Organization Reference Implementation

Phase	What Happened	Key Learning
1: Awareness	The dual-path fabrication incident surfaced during active research — a plausible result propagated through a relay chain without detection. This served as the unplanned awareness	Awareness events often come from visible failures. The goal of Phase 1 is to create structured awareness before the failure occurs.

Phase	What Happened	Key Learning
	event.	
2: Assessment	Retrospective mapping identified five high-risk entry points: PSLQ output, relay chain handoffs, multi-session trust inheritance, scan result citation, and constant catalog entries.	The Epistemic Risk Map revealed that the highest-risk entry point (relay chains) had no verification step at all — a trust gap of 3 (T0 treated as T3).
3: Pilot	ZTEK minimum viable implementation applied to one scan workflow: reproduce commands, evidence classes, boundary verification. The pilot caught three unverified claims in eight weeks that would previously have entered the catalog.	The overhead was approximately 20% additional time per scan cycle. The team accepted this as worthwhile after the first caught fabrication — a claim already referenced in a draft paper section.
4: Institutionalization	ZTEK pillars encoded into research workflow protocol. Trust tier requirements added to paper submission checklist. AI disclosure statements standardized across all papers in the series.	Incentive alignment was the hardest change: the natural incentive in research is to advance. Making verification a named contribution — a concrete research output, not overhead — helped significantly.
5: Continuous Improvement	Quarterly audit process established. Failure mode registry maintained with all five documented incidents. Model upgrade protocol applied when switching between Claude versions.	The failure mode registry became the most valuable artifact — a living document that directly grounded Papers 9 and 10 of this series with empirically derived protocol pillars.

14. Pilot Vignette: Before and After Metrics

The following vignette describes the ZTEK pilot conducted within the PCF mathematical research program. It is provided as a concrete reference for organizations designing their own Phase 3 pilots.

14.1 Context

The pilot workflow was a numerical constant verification scan: a repeating cycle in which candidate polynomial continued fraction constants are generated by AI (GitHub Copilot / Claude), assessed for irrationality, and entered into the research catalog. The scan cycle runs approximately weekly and produces 5–15 candidate constants per run. (Source: PCF research program, 2024–2025 pilot.)

The workflow was selected because: (a) it had clear, verifiable outputs (numerical constants at specified precision), (b) it was entirely AI-assisted with no existing verification gate, and (c) a fabricated constant entering the catalog would require significant rework to detect and remove.

14.2 Baseline Run (8 weeks, without ZTEK)

Metric	Observed Value	Notes
Total candidate constants generated	47	Across 8 weekly scan cycles
Constants entered into catalog	31	66% of generated; remainder discarded by researcher judgment
Provenance records available	0 of 31	No reproduce commands, no environment fingerprints
Independent verification performed	3 of 31	Ad hoc; no systematic process
Fabricated or trivially incorrect claims	Estimated 4–6	Identified retrospectively during ZTEK implementation
Catalog corrections required post-scan	2 confirmed	Discovered only when referenced in draft paper

14.3 Brake-Enabled Run (8 weeks, with ZTEK minimum viable implementation)

Metric	Observed Value	Change vs. Baseline
Total candidate constants generated	41	Comparable volume
Constants entering catalog (T1 or above)	12	Down from 31 — gate is working
Provenance records completed	12 of 12	100% — required for gate passage
Independent verification performed	9 of 12	75% — systematic, not ad hoc
Fabricated claims caught at gate	3 confirmed	0 in baseline — systematic detection active
Catalog corrections required post-scan	0	No downstream corrections needed in pilot period
Verification overhead per scan cycle	~20-25% additional time	Accepted as worthwhile by research team

14.4 Interpretation

The most significant finding is not the fabrication catch rate — it is the provenance coverage shift from 0% to 100%. This single change means that every claim now entering the catalog has a reproduce command attached. Future researchers (including the original researcher in a new session) can verify any catalog entry independently, without trusting session memory or narrative context.

The reduction in catalog entries (31 to 12) reflects the gate working as designed: claims that cannot prove provenance do not enter. This is not a loss of productivity — it is a precision gain. Twelve verified entries are more valuable than thirty-one unverified ones.

15. Adoption Roadmap

Timeline	Phase	Key Activities	Success Gate
Weeks 1–2	Phase 1: Awareness	Literacy brief for leadership; fabrication assumption exercise; visible failure audit	Leadership commits to assessment
Weeks 3–4	Phase 2: Assessment	Workflow mapping; trust tier audit; relay chain identification; Epistemic Risk Map	Risk Map complete; pilot workflow selected
Weeks 5–12	Phase 3: Pilot	Baseline run (4 weeks); brake-enabled run (4 weeks); metrics comparison	Measurable fabrication reduction demonstrated
Weeks 13–20	Phase 4: Institutionalization	Policy drafting; process embedding; incentive changes; training	Policy published; 3+ workflows compliant
Ongoing (quarterly)	Phase 5: Continuous Improvement	Quarterly audits; failure mode registry; model upgrade protocol	Audit completion; registry maintained

16. Conclusion

The AI Epistemic Governance Stack is complete. Paper 7 explains why AI fails. Paper 8 tracks the failure. Paper 9 enforces the boundary. Paper 10 moves the framework from individual practice to organizational norm.

The change management framework proposed here is deliberately pragmatic. It does not require organizational transformation before it produces value. The Epistemic Risk Map alone — a single Phase 2 deliverable — gives any organization an immediately actionable picture of its AI-related epistemic exposure.

The deeper argument of this paper is this: AI governance is not primarily a technical problem. The brakes work — the pilot vignette demonstrates this with concrete metrics. The harder challenge is building the organizational culture, incentive structures, and verification disciplines that ensure the brakes are actually used. That is a change management problem — and like all change management problems, it is solved not by mandate but by making the right behavior

visible, rewarded, and easier than the alternative.

Although the reference implementation in this paper is drawn from a mathematical research program, the framework applies to any organizational context where AI-generated output informs decisions, documents, or downstream systems — including legal, finance, strategy, product development, and public communications teams. The Epistemic Risk Map construction in Section 12 requires no domain-specific knowledge and can be completed by any team in a single working session.

The organizations that solve this problem first will not just produce more reliable AI-assisted work. They will build a durable institutional capability — the ability to integrate AI into their knowledge workflows without sacrificing the epistemic standards that make those workflows trustworthy. The Epistemic Risk Map is the recommended starting point for any organization: a single deliverable, no technical prerequisites, immediate value.

References

[1] AI Behavior Science — Founding Territory Paper. Zenodo. DOI: 10.5281/zenodo.19562751

[2] AEAL: Agent Epistemic Accountability Layer. Zenodo. DOI: 10.5281/zenodo.19565086

[3] ZTEK: A Zero Trust Epistemic Kernel for AI-Assisted Research Workflows. [This series, Paper 9]

[4] Kotter, J.P. Leading Change. Harvard Business Review Press. 1996.

[5] Prosci. ADKAR: A Model for Change in Business, Government and our Community. 2006.

[6] Rose, J. et al. Zero Trust Architecture. NIST Special Publication 800-207. 2020.

[7] Gioia, D.A. and Chittipeddi, K. Sensemaking and Sensegiving in Strategic Change Initiation. Strategic Management Journal, 12(6), 433–448. 1991.

[8] Ji, Z. et al. Survey of Hallucination in Natural Language Generation. ACM Computing Surveys, 55(12), 1–38. 2023. https://doi.org/10.1145/3571730

AI Disclosure:
This paper was developed in collaboration with Claude (Anthropic) as a strategic research partner. The five-phase framework, resistance map, role structure, and pilot vignette were developed through iterative dialogue grounded in the author's direct organizational and research experience. All framework design decisions were made and verified by the human author.